

An introduction to numerical methods using BLAS

Georgios Kafanas

GreekLUG

26 November 2023

Overview

- What are numerical libraries and how to use them?
- Functionality and uses of BLAS
- How computer architecture (caches) affect implementation of computations

Not all evaluations are the same!

Consider a simple matrix multiplication

$$\begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} \cdot \begin{pmatrix} 1 \\ 2 \end{pmatrix}. \quad (1)$$

Can be evaluated in 2 ways, with a dot product or a scalar-matrix product:

$$\begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} \cdot \begin{pmatrix} 1 \\ 2 \end{pmatrix} = \begin{pmatrix} \begin{pmatrix} 1 & 2 \end{pmatrix} \cdot \begin{pmatrix} 1 \\ 2 \end{pmatrix} \\ \begin{pmatrix} 3 & 4 \end{pmatrix} \cdot \begin{pmatrix} 1 \\ 2 \end{pmatrix} \end{pmatrix}, \quad \begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} \cdot \begin{pmatrix} 1 \\ 2 \end{pmatrix} = 1 \begin{pmatrix} 1 \\ 3 \end{pmatrix} + 2 \begin{pmatrix} 2 \\ 4 \end{pmatrix}.$$

Software libraries

- BLAS is a typical software library
- Libraries can be used in 2 forms:
 - Static
 - Dynamic


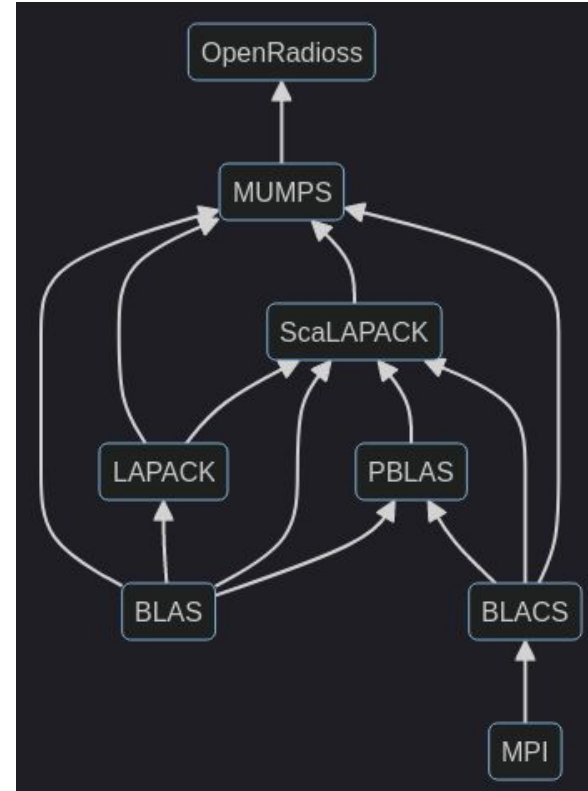
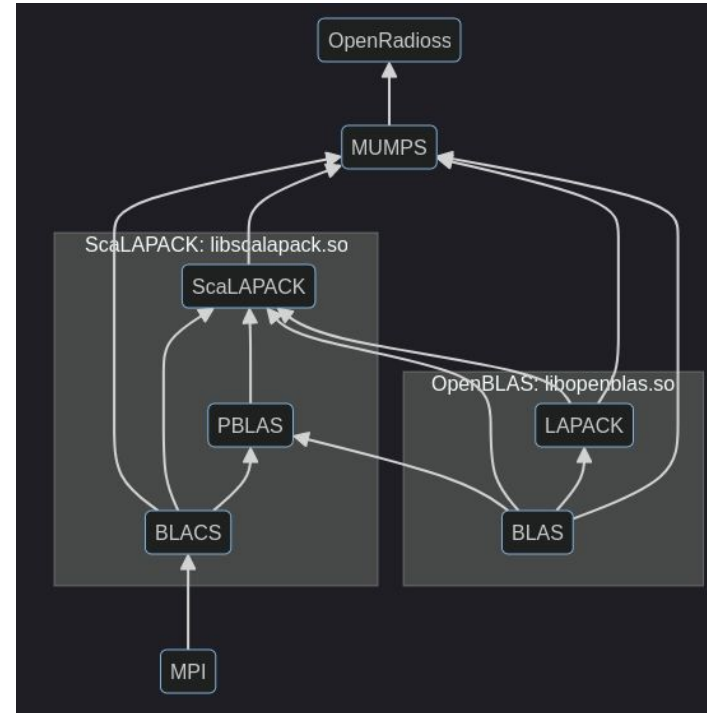
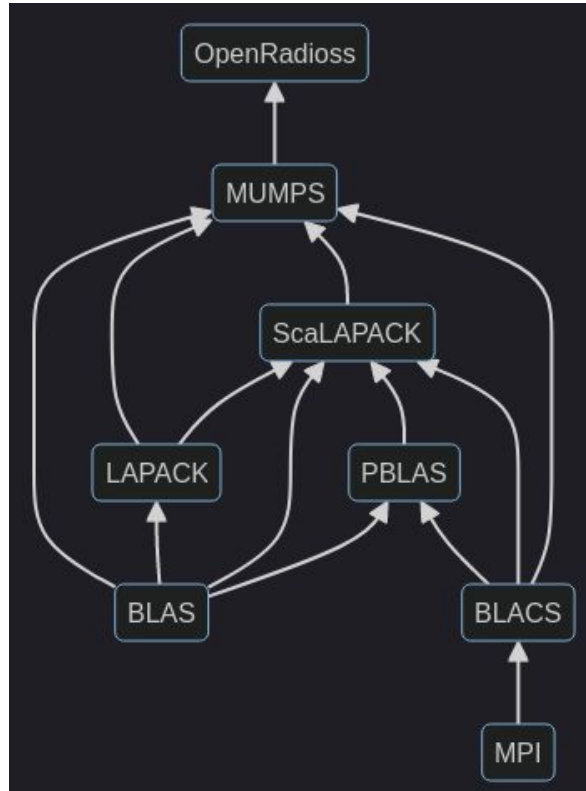


Figure 1: Shared library dependencies of an example application.



Software libraries



Practical session

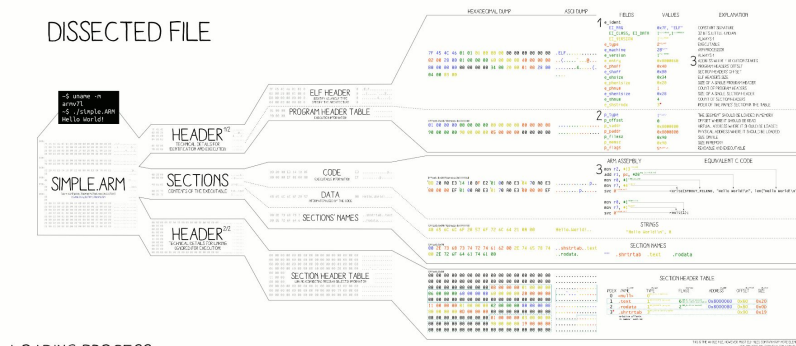
Compile and install:

- BLAS: <https://gitlab.com/greeklug/lapack/-/tree/greeklug-presentation>
- Matrix-Market I/O library:
https://gitlab.com/greeklug/matrix_market_exchange_formats

Tools for inspecting libraries and executables

- Does this `libopeblas.so` instance implement the CBLAS interface?


ELF¹⁰¹ a Linux executable walk-through ANGE ALBERTINI CORKAMP.COM



LOADING PROCESS

1 HEADER

THE ELF HEADER IS PARSED
THE PROGRAM HEADER IS PARSED
(SECTIONS ARE NOT USED)



2 MAPPING

THE FILE IS MAPPED IN MEMORY
ACCORDING TO ITS SEGMENTS

3 EXECUTION

ENTRY IS CALLED
SYSCALLS² ARE ACCESSED VIA:
- SYSCALL NUMBER IN THE R7 REGISTER
- CALLING INSTRUCTION SVC

TRIVIA

THE ELF WAS FIRST SPECIFIED BY U.S. "C" AND "IT"
FOR UNIX SYSTEM V, IN 1989

THE ELF IS USED, AMONG OTHERS, IN:
- LINUX, ANDROID, BSD, SOLARIS, BEOS
- PSP, PLAYSTATION 2-4, DREAMCAST, GAMECUBE, Wii
- VARIOUS OSes MADE BY SAMSUNG, ERICSSON, NOKIA
- MICROCONTROLLERS FROM ATMEL, TEXAS INSTRUMENTS

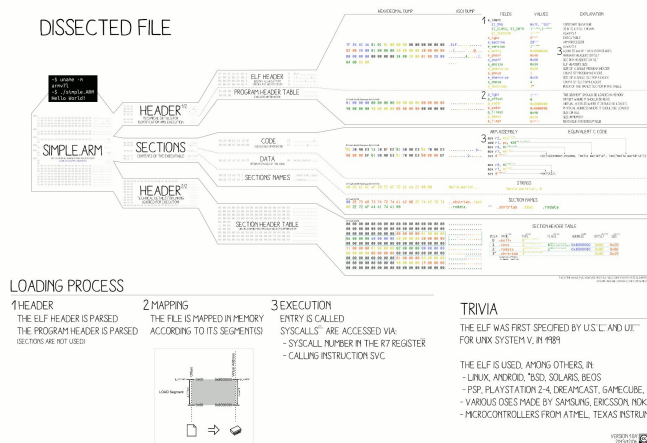
VERSION 1001
2020/07/24

Tools for inspecting libraries and executables

To investigate the shared object:

- **readelf**: display information about ELF files
 - **--all**: all sections
 - **--file-header**: information about interoperability
 - **--dynamic**: dynamically linked libraries and other information
- **objdump**: display information about objects
 - **--syms**: information for symbols (functions and variables)
 - **--demangle**: restore human readable names for objects generated from C++
- **nm**: list symbols
 - **--dynamic**: list only export symbols (only for dynamic libraries)

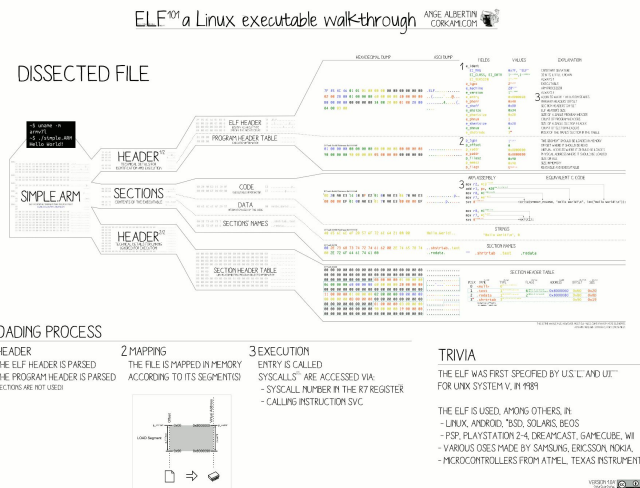
ELFTM a Linux executable walk-through ANGE ALBERTO CORRAVECO



Tools for inspecting libraries and executables

Even extract information about function signatures (needs debug info, $-g$):

- **Read debug info with** `readelf`
 - `--debug-dump=info`
- **Partially disassemble with** `objdump`
 - `--disassemble`
 - `--disassemble-all`



Practical session

- Compile the tutorial example code: <https://gitlab.com/greeklug/blas-tutorial>
- Call some function Matrix Market I/O
- Can you break the linking? Try removing the linker option: `--no-as-needed`





Figure 2: Wrong symbol resolution after relinking the example application.

Data representation

- Computer memory is linear
- Matrices are linearized:




```
typedef struct _dense_matrix {  
    /* Data structure storing matrix A */  
    double* a; // Pointer to the C array with the entries of A  
    int m;     // Number of rows in A  
    int n;     // Number of columns in A  
} dense_matrix;
```

Caching

- Direct linearization is not sufficient for good performance!
- Caches affect the speed of memory access

```
for ( int i = 0; i < n; ++i ) {  
    a[i] = 0;  
}
```



*up to 64 for some special SIMD instructions sets such as AVX-512

Caching

- Direct linearization is not sufficient for good performance!
- Caches affect the speed of memory access

```
for ( int i = 0; i < n; ++i )  
{  
    a[i] = 0;  
}
```

Vectorizable:



Non-vectorizable:




Caching

- Direct linearization is not sufficient for good performance!
- Caches affect the speed of memory access

```
#pragma omp simd aligned(a:32)
for ( int i = 0; i < 4*n; i+=1 ) {
    a[i] = 0;
}
```

Vectorizable:



Non-vectorizable:




Caching

- Direct linearization is not sufficient for good performance!
- Caches affect the speed of memory access

```
for ( int i = 0; i < n; i+=1 ) {  
    a[4*i] = 0;  
    a[4*i+1] = 0;  
    a[4*i+2] = 0;  
    a[4*i+3] = 0;  
}
```

Vectorizable:



Non-vectorizable:



Not all evaluations are the same!

Revisiting the simple matrix-vector product:

$$\begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} \cdot \begin{pmatrix} 1 \\ 2 \end{pmatrix}. \quad (1)$$

The dot product evaluation jumps across cache lines!

$$\begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} \cdot \begin{pmatrix} 1 \\ 2 \end{pmatrix} = \begin{pmatrix} \begin{pmatrix} 1 & 2 \end{pmatrix} \cdot \begin{pmatrix} 1 \\ 2 \end{pmatrix} \\ \begin{pmatrix} 3 & 4 \end{pmatrix} \cdot \begin{pmatrix} 1 \\ 2 \end{pmatrix} \end{pmatrix}, \quad \begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} \cdot \begin{pmatrix} 1 \\ 2 \end{pmatrix} = 1 \begin{pmatrix} 1 \\ 3 \end{pmatrix} + 2 \begin{pmatrix} 2 \\ 4 \end{pmatrix}.$$

Caching

- Direct linearization is not sufficient for good performance!
- Caches affect the speed of memory access

```
typedef struct _dense_matrix {  
    /* Data structure storing matrix A */  
    double* a; // Pointer to the C array with the entries of A  
    int m;     // Number of rows in A  
    int n;     // Number of columns in A  
    int nzmax; // Maximum number of entries that can be stored in array a  
    int lda;   // Leading dimension of the array A  
} dense_matrix;
```

Practical session

- Call some function (DGEMV) of BLAS
- Try the code with aligned memory allocation!

BLAS naming conventions

- Operations organized by computational complexity
 - Level 1: $O(n)$
 - Level 2: $O(n^2)$
 - Level 3: $O(n^3)$
- BLAS supports various number types and numerical precision (first part of function names):
 - single precision (**S**) with 32-bits,
 - double precision (**D**) with 64-bits,
 - single precision complex (**C**) with 64-bits, and
 - double precision complex (**Z**) with 128-bits.

BLAS naming conventions

- Matrix properties are exploited to save space and reduce memory accesses:

GE: $\begin{pmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{pmatrix} \longrightarrow \begin{array}{|c|c|c|c|c|c|c|c|c|} \hline a_{11} & a_{21} & a_{31} & a_{12} & a_{22} & a_{32} & a_{13} & a_{23} & a_{33} \\ \hline \end{array}$

SY: $\begin{pmatrix} a_{11} & a_{12} & a_{13} \\ a_{12} & a_{22} & a_{23} \\ a_{13} & a_{23} & a_{33} \end{pmatrix} \longrightarrow \begin{array}{|c|c|c|c|c|c|c|c|c|} \hline a_{11} & a_{21} & a_{31} & a_{12} & a_{22} & * & a_{13} & * & * \\ \hline \end{array}$

TR: $\begin{pmatrix} a_{11} & a_{12} & a_{13} \\ 0 & a_{22} & a_{23} \\ 0 & 0 & a_{33} \end{pmatrix} \longrightarrow \begin{array}{|c|c|c|c|c|c|c|c|c|} \hline a_{11} & a_{21} & a_{31} & a_{12} & a_{22} & * & a_{13} & * & * \\ \hline \end{array}$

SP: $\begin{pmatrix} a_{11} & a_{12} & a_{13} \\ a_{12} & a_{22} & a_{23} \\ a_{13} & a_{23} & a_{33} \end{pmatrix} \longrightarrow \begin{array}{|c|c|c|c|c|c|} \hline a_{11} & a_{21} & a_{31} & a_{12} & a_{22} & a_{13} \\ \hline \end{array}$

TP: $\begin{pmatrix} a_{11} & a_{12} & a_{13} \\ 0 & a_{22} & a_{23} \\ 0 & 0 & a_{33} \end{pmatrix} \longrightarrow \begin{array}{|c|c|c|c|c|c|} \hline a_{11} & a_{21} & a_{31} & a_{12} & a_{22} & a_{13} \\ \hline \end{array}$

BLAS naming conventions

- Matrix properties are exploited to save space and reduce memory accesses.
- This forms the second part of the name:

	Storage type		
Algebraic properties	Standard (-)	Banded (B)	Packed (P)
General (G)	GE	GB	
Symmetric (S)	SY	SB	SP
Hermitian (H)	HE	HB	HP
Triangular (T)	TR	TB	TP

BLAS naming conventions

- Last part is the type of the operands:
 - V: vector
 - M: matrix
- For instance:

DGEMV:

- D: double precision
- GE: general matrix
- MV: matrix-vector multiplication

DGEMM:

- D: double precision
- GE: general matrix
- MM: matrix-matrix multiplication

- The convention does not work always, especially for Level 1 operations
 - DAXPY: $y \leftarrow ax+y$

Course notes and resources

Will appear in the tutorial directory: <https://gitlab.com/greeklug/blas-tutorial>

Official BLAS webpage: <https://www.netlib.org/blas/>

- Quick reference (function list): <https://www.netlib.org/blas/>
- Reference BLAS implementation:
https://www.netlib.org/lapack/explore-html/d1/df9/group__blas.html

Seminar on numerical methods coming soon in EuroCC:
<https://www.eurocc-access.eu/services/training/>



ABOUT US ▾

SERVICES ▾

NEWS ▾

LOGIN

Training



Find upcoming training offers from the National Competence Centres here!

From beginner to expert courses, whether C++ or Open MP, there's something for everybody. If you want to see all training offers, including past courses, go to the [HPC in Europe Portal](#)!

Event feed:

Download

Copy URL

Search events:

Type text to search for

Difficulty Level

Language of
Instruction

Country

Audience

Format

Scientific Domain

Technical Domain

Project



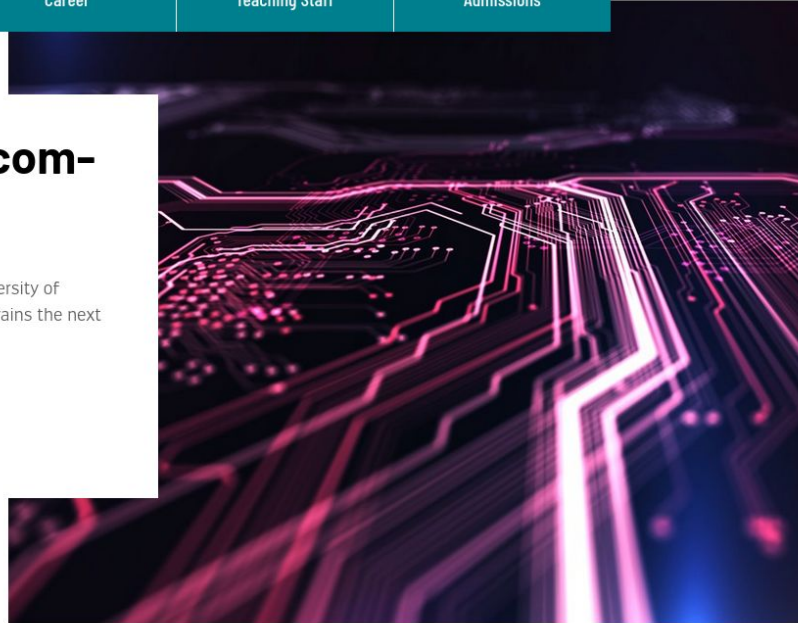
Thank you!

Deepen your career in computing

The Master in High Performance Computing (HPC) at the University of Luxembourg is a 2-year innovative Master's programme that trains the next generation of HPC experts in Luxembourg and Europe.

[Learn more](#)

[Teachers](#)



The programme at a glance – 120 ECTS

 Duration:
2 years / 4 sem


 Language:
EN

 Admissions
EU: 1 Feb-28 Aug 2024
Non-EU: 1 Feb-30 Apr 2024

 Fees:
€200/sem

 Format: **Full time**

 Campus: **Belval**

 Available places: **40**

Launch your career in Mathematical Modelling and Computational Sciences

The two-year track in Mathematical Modelling and Computational Sciences of the Master in Mathematics focuses on both computational and fundamental aspects of mathematics and prepares for immediate employment after graduation. You will become thoroughly acquainted with industrial mathematics due to the possibility of doing a summer internship and your master's thesis with a local company and thus require fruitful insights into the real world.

[Learn more](#)

[Teachers](#)

A woman with dark hair and glasses, wearing a light blue and white striped shirt, stands next to a large digital screen. She is gesturing with her right hand while holding a tablet in her left. The screen displays several colorful dice (purple, green, silver, pink) and the text "choisit un des deux dés restants".

choisit un des deux dés restants

The Programme at a glance – 120 ECTS



Duration:
2 years / 4 sem



Teaching Languages:
EN



Admissions:
EU: 1 Feb 2023 – 31 Aug
2023
Non-EU: 1 Feb 2023 –
30 Apr 2023



Fees:
200€/ sem. (semester
1)

EUMaster4HPC

European Master
For High Performance Computing

Application

Summer School

Research & Industry

Applicants from participant countries are eligible for receiving EU funds.



- Master in High Performance Computing:
<https://www.uni.lu/fstm-en/study-programs/master-in-high-performance-computing/>
- Master in Mathematics - Mathematical Modelling and Computational Sciences:
<https://www.uni.lu/fstm-en/study-programs/master-in-mathematics-mathematical-modelling-and-computational-sciences/>
- European Master For High Performance Computing:
<https://eumaster4hpc.uni.lu/>